

## Case Study 1: Estimating Click Probabilities

### SGD cont'd AdaGrad

Machine Learning for Big Data  
CSE547/STAT548, University of Washington

Emily Fox

April 2<sup>nd</sup>, 2015

©Emily Fox 2015

1

## Learning Problem for Click Prediction

- Prediction task:  $X \rightarrow \{0, 1\}$        $P(\text{click}=1 | X)$
- Features:  $X = (\text{feats of page, ad, user})$
- Data:  $(x^i, y^i)$        $(\text{webpage1, ad7, user25, time12}) \leftarrow x^i$   
 $\text{click}=1 \leftarrow y^i$ 
  - Batch:  $(x^1, y^1), \dots, (x^n, y^n)$
  - Online:  $\text{data as a stream}$   
 $\text{user arrives at a page} \rightarrow X^t$        $\begin{matrix} \text{predict } y \\ \text{click?} \end{matrix}$   
 $\text{observe } y^t$
- Many approaches (e.g., logistic regression, SVMs, naïve Bayes, decision trees, boosting,...)
  - Focus on logistic regression; captures main concepts, ideas generalize to other approaches

©Emily Fox 2015

2

# Standard v. Regularized Updates

- Maximum conditional likelihood estimate

$$\mathbf{w}^* = \arg \max_{\mathbf{w}} \ln \left[ \prod_{j=1}^N P(y^j | \mathbf{x}^j, \mathbf{w}) \right]$$

$$w_i^{(t+1)} \leftarrow w_i^{(t)} + \eta \sum_j x_i^j [y^j - \hat{P}(Y^j = 1 | \mathbf{x}^j, \mathbf{w}^{(t)})]$$

- Regularized maximum conditional likelihood estimate

$$\mathbf{w}^* = \arg \max_{\mathbf{w}} \ln \left[ \prod_j P(y^j | \mathbf{x}^j, \mathbf{w}) \right] - \frac{\lambda}{2} \sum_{i>0} w_i^2$$

$$w_i^{(t+1)} \leftarrow w_i^{(t)} + \eta \left\{ -\lambda w_i^{(t)} + \sum_j x_i^j [y^j - \hat{P}(Y^j = 1 | \mathbf{x}^j, \mathbf{w}^{(t)})] \right\}$$

*neg. derivative ← more towards 0*

©Emily Fox 2015

3

## Challenge 1: Complexity of computing gradients $\partial$ features

- What's the cost of a gradient update step for LR???

$$w_i^{(t+1)} \leftarrow w_i^{(t)} + \eta \left\{ -\lambda w_i^{(t)} + \sum_{j=1}^N x_i^j [y^j - \hat{P}(Y^j = 1 | \mathbf{x}^j, \mathbf{w}^{(t)})] \right\}$$

*for each  $i$*

*$O(d)$*

*$w_i^{(t+1)}$*

*$O(Nd)$*

*$\forall$  features  $i$ , cost is  $O(Nd^2)$  ... can cache  $p(y^j=1 | \mathbf{x}^j, \mathbf{w}^{(t)})$*

*$O(Nd)$*

*In "big data" -  $N$  is very large  
 $O(Nd)$  for only taking little  $\eta$  step*

©Emily Fox 2015

4

## Challenge 2: Data is streaming

- Assumption thus far: Batch data
- But, click prediction is a streaming data task:

- User enters query, and ad must be selected:
- Observe  $x^j$ , and must predict  $\hat{y}^j$

$x \rightarrow \text{[ ]} \rightarrow x^j \rightarrow \text{predict } \hat{y}^j \rightarrow \text{show ad}$

- User either clicks or doesn't click on ad:
- Label  $y^j$  is revealed afterwards
- Google gets a reward if user clicks on ad

- Weights must be updated for next time:

$$w^{(t+1)} \leftarrow w^{(t)} + \Delta \quad \text{depends just on recent example(s)}$$

©Emily Fox 2015

5

## SGD: Stochastic Gradient Ascent (or Descent)

- “True” gradient:  $\nabla \ell(\mathbf{w}) = E_{\mathbf{x}} [\nabla \ell(\mathbf{w}, \mathbf{x})]$
- Sample based approximation:  $x^j \stackrel{iid}{\sim} p(x)$   
 $\nabla \ell(\mathbf{w}) = E_{\mathbf{x}} [\nabla \ell(\mathbf{w}, \mathbf{x})] \approx \hat{\nabla} \ell(\mathbf{w}) = \frac{1}{N} \sum_{j=1}^N \nabla \ell(\mathbf{w}, x^j)$   
 the bigger  $N$ , the closer  $\hat{\nabla} \ell$  to  $\nabla \ell$
- What if we estimate gradient with just one sample???  $N=1$ 
  - Unbiased estimate of gradient
  - Very noisy!
  - Called stochastic gradient ascent (or descent)
  - Among many other names
  - VERY useful in practice!!!

©Emily Fox 2015

6

## Stochastic Gradient Ascent: General Case

- Given a stochastic function of parameters:
  - Want to find maximum
- Start from  $\mathbf{w}^{(0)}$
- Repeat until convergence:
  - Get a sample data point  $\mathbf{x}^t$
  - Update parameters:
- Works in the online learning setting!
- Complexity of each gradient step is constant in number of examples!
- In general, step size changes with iterations

©Emily Fox 2015

7

## Stochastic Gradient Ascent for Logistic Regression

- Logistic loss as a stochastic function:

$$E_{\mathbf{x}} [\ell(\mathbf{w}, \mathbf{x})] = E_{\mathbf{x}} \left[ \ln P(y|\mathbf{x}, \mathbf{w}) - \frac{\lambda}{2} \|\mathbf{w}\|_2^2 \right]$$

- Batch gradient ascent updates:

$$w_i^{(t+1)} \leftarrow w_i^{(t)} + \eta \left\{ -\lambda w_i^{(t)} + \frac{1}{N} \sum_{j=1}^N x_i^{(j)} [y^{(j)} - P(Y = 1 | \mathbf{x}^{(j)}, \mathbf{w}^{(t)})] \right\}$$

- Stochastic gradient ascent updates:

- Online setting:

$$w_i^{(t+1)} \leftarrow w_i^{(t)} + \eta_t \left\{ -\lambda w_i^{(t)} + x_i^{(t)} [y^{(t)} - P(Y = 1 | \mathbf{x}^{(t)}, \mathbf{w}^{(t)})] \right\}$$

©Emily Fox 2015

8

# Convergence Rate of SGD

- **Theorem:**
  - (see Nemirovski et al '09 from readings)
  - Let  $\ell$  be a strongly convex stochastic function
  - Assume gradient of  $\ell$  is Lipschitz continuous and bounded
  - Then, for step sizes:
    - The expected loss decreases as  $O(1/t)$ :

©Emily Fox 2015

9

# Convergence Rates for Gradient Descent/Ascent vs. SGD

- Number of Iterations to get to accuracy
$$\ell(\mathbf{w}^*) - \ell(\mathbf{w}) \leq \epsilon$$
- Gradient descent:
  - If func is strongly convex:  $O(\ln(1/\epsilon))$  iterations
- Stochastic gradient descent:
  - If func is strongly convex:  $O(1/\epsilon)$  iterations
- Seems exponentially worse, but much more subtle:
  - Total running time, e.g., for logistic regression:
    - Gradient descent:
    - SGD:
    - SGD can win when we have a lot of data
  - See readings for more details

©Emily Fox 2015

10

## Constrained SGD: Projected Gradient

- Consider an arbitrary restricted feature space  $\mathbf{w} \in \mathcal{W}$
- Optimization objective:
- If  $\mathbf{w} \in \mathcal{W}$ , can use **projected gradient** for (sub)gradient descent

$$\mathbf{w}^{(t+1)} =$$

©Emily Fox 2015

11


## Motivating AdaGrad (Duchi, Hazan, Singer 2011)

- Assuming  $\mathbf{w} \in \mathbb{R}^d$ , standard stochastic (sub)gradient descent updates are of the form:
$$w_i^{(t+1)} \leftarrow w_i^{(t)} - \eta_t g_{t,i}$$
- Should all features share the same learning rate?
- Often have high-dimensional feature spaces
  - Many features are irrelevant
  - Rare features are often very informative
- Adagrad provides a feature-specific adaptive learning rate by incorporating knowledge of the geometry of past observations

©Emily Fox 2015

12

## Why Adapt to Geometry?



$y_t$	$x_{t,1}$	$x_{t,2}$	$x_{t,3}$
1	1	0	0
-1	.5	0	1
1	-.5	1	0
-1	0	0	0
1	.5	0	0
-1	1	0	0
1	-.1	1	0
-1	-.5	0	1

Examples from  
Duchi et al.  
ISMP 2012  
slides

- ① Frequent, irrelevant
- ② Infrequent, predictive
- ③ Infrequent, predictive

©Emily Fox 2015

13

## Not All Features are Created Equal

- Examples:

Text data:

The most unsung birthday  
in American business and  
technological history  
this year may be the 50th  
anniversary of the Xerox  
914 photocopier.<sup>a</sup>

<sup>a</sup>The Atlantic, July/August 2010.

High-dimensional image features




Images from Duchi et al. ISMP 2012 slides

©Emily Fox 2015

14

## Visualizing Effect



Credit:  
<http://imgur.com/a/Hqolp>

©Emily Fox 2015

15

## Regret Minimization

- How do we assess the performance of an online algorithm?
- Algorithm iteratively predicts  $\mathbf{w}^{(t)}$
- Incur **loss**  $\ell_t(\mathbf{w}^{(t)})$
- **Regret:**  
What is the total incurred loss of algorithm relative to the best choice of  $\mathbf{W}$  that could have been made **retrospectively**

$$R(T) = \sum_{t=1}^T \ell_t(\mathbf{w}^{(t)}) - \inf_{\mathbf{w} \in \mathcal{W}} \sum_{t=1}^T \ell_t(\mathbf{w})$$

©Emily Fox 2015

16

## Regret Bounds for Standard SGD

- Standard projected gradient stochastic updates:

$$\mathbf{w}^{(t+1)} = \arg \min_{\mathbf{w} \in \mathcal{W}} \|\mathbf{w} - (\mathbf{w}^{(t)} - \eta g_t)\|_2^2$$

- Standard regret bound:

$$\sum_{t=1}^T \ell_t(\mathbf{w}^{(t)}) - \ell_t(\mathbf{w}^*) \leq \frac{1}{2\eta} \|\mathbf{w}^{(1)} - \mathbf{w}^*\|_2^2 + \frac{\eta}{2} \sum_{t=1}^T \|g_t\|_2^2$$

©Emily Fox 2015

17

## Projected Gradient using Mahalanobis

- Standard projected gradient stochastic updates:

$$\mathbf{w}^{(t+1)} = \arg \min_{\mathbf{w} \in \mathcal{W}} \|\mathbf{w} - (\mathbf{w}^{(t)} - \eta g_t)\|_2^2$$

- What if instead of an  $L_2$  metric for projection, we considered the **Mahalanobis** norm

$$\mathbf{w}^{(t+1)} = \arg \min_{\mathbf{w} \in \mathcal{W}} \|\mathbf{w} - (\mathbf{w}^{(t)} - \eta A^{-1} g_t)\|_A^2$$

©Emily Fox 2015

18

## Mahalanobis Regret Bounds

$$\mathbf{w}^{(t+1)} = \arg \min_{\mathbf{w} \in \mathcal{W}} \|\mathbf{w} - (\mathbf{w}^{(t)} - \eta A^{-1} g_t)\|_A^2$$

- **What  $A$  to choose?**
- Regret bound now:

$$\sum_{t=1}^T \ell_t(\mathbf{w}^{(t)}) - \ell_t(\mathbf{w}^*) \leq \frac{1}{2\eta} \|\mathbf{w}^{(1)} - \mathbf{w}^*\|_2^2 + \frac{\eta}{2} \sum_{t=1}^T \|g_t\|_{A^{-1}}^2$$

- What if we minimize upper bound on regret w.r.t.  $A$  in hindsight?

$$\min_A \sum_{t=1}^T g_t^T A^{-1} g_t$$

©Emily Fox 2015

19

## Mahalanobis Regret Minimization

- Objective:

$$\min_A \sum_{t=1}^T g_t^T A^{-1} g_t \quad \text{subject to } A \succeq 0, \text{tr}(A) \leq C$$

- Solution:

$$A = c \left( \sum_{t=1}^T g_t g_t^T \right)^{\frac{1}{2}}$$

For proof, see Appendix E, Lemma 15 of Duchi et al. 2011.  
Uses “trace trick” and Lagrangian.

- $A$  defines the norm of the metric space we should be operating in

©Emily Fox 2015

20

## AdaGrad Algorithm

- At time  $t$ , estimate optimal (sub)gradient modification  $A$  by

$$A_t = \left( \sum_{\tau=1}^t g_\tau g_\tau^T \right)^{\frac{1}{2}}$$

- For  $d$  large,  $A_t$  is computationally intensive to compute. Instead,

- Then, algorithm is a simple modification of normal updates:

$$\mathbf{w}^{(t+1)} = \arg \min_{\mathbf{w} \in \mathcal{W}} \|\mathbf{w} - (\mathbf{w}^{(t)} - \eta \text{diag}(A_t)^{-1} g_t)\|_{\text{diag}(A_t)}^2$$

©Emily Fox 2015

21

## AdaGrad in Euclidean Space

- For  $\mathcal{W} = \mathbb{R}^d$ ,
- For each feature dimension,

$$w_i^{(t+1)} \leftarrow w_i^{(t)} - \eta_{t,i} g_{t,i}$$

where

$$\eta_{t,i} =$$

- That is,

$$w_i^{(t+1)} \leftarrow w_i^{(t)} - \frac{\eta}{\sqrt{\sum_{\tau=1}^t g_{\tau,i}^2}} g_{t,i}$$

- Each feature dimension has its own learning rate!
  - Adapts with  $t$
  - Takes geometry of the past observations into account
  - Primary role of  $\eta$  is determining rate the first time a feature is encountered

©Emily Fox 2015

22

## AdaGrad Theoretical Guarantees

- AdaGrad regret bound:

$$\sum_{t=1}^T \ell_t(\mathbf{w}^{(t)}) - \ell_t(\mathbf{w}^*) \leq 2R_\infty \sum_{i=1}^d \|g_{1:T,i}\|_2$$

$R_\infty := \max_t \|\mathbf{w}^{(t)} - \mathbf{w}^*\|_\infty$

- In stochastic setting:

$$\mathbb{E} \left[ \ell \left( \frac{1}{T} \sum_{t=1}^T \mathbf{w}^{(t)} \right) \right] - \ell(\mathbf{w}^*) \leq \frac{2R_\infty}{T} \sum_{i=1}^d \mathbb{E}[\|g_{1:T,i}\|_2]$$

- This really is used in practice!
- Many cool examples. Let's just examine one...

©Emily Fox 2015

23

## AdaGrad Theoretical Example

- Expect to out-perform when gradient vectors are *sparse*
- SVM hinge loss example:

$$\ell_t(\mathbf{w}) = [1 - y^t \langle \mathbf{x}^t, \mathbf{w} \rangle]_+$$

$$\mathbf{x}^t \in \{-1, 0, 1\}^d$$

- If  $x_j^t \neq 0$  with probability  $\propto j^{-\alpha}$ ,  $\alpha > 1$

$$\mathbb{E} \left[ \ell \left( \frac{1}{T} \sum_{t=1}^T \mathbf{w}^{(t)} \right) \right] - \ell(\mathbf{w}^*) = \mathcal{O} \left( \frac{\|\mathbf{w}^*\|_\infty}{\sqrt{T}} \cdot \max\{\log d, d^{1-\alpha/2}\} \right)$$

- Previously best known method:  $\mathbb{E} \left[ \ell \left( \frac{1}{T} \sum_{t=1}^T \mathbf{w}^{(t)} \right) \right] - \ell(\mathbf{w}^*) = \mathcal{O} \left( \frac{\|\mathbf{w}^*\|_\infty}{\sqrt{T}} \cdot \sqrt{d} \right)$


©Emily Fox 2015

24

# Neural Network Learning

- Very non-convex problem, but use SGD methods anyway

$$\ell(w, x) = \log(1 + \exp(\langle [p(\langle w_1, x_1 \rangle) \cdots p(\langle w_k, x_k \rangle)], x_0 \rangle))$$



(Dean et al. 2012)

Distributed,  $d = 1.7 \cdot 10^9$  parameters. SGD and AdaGrad use 80 machines (1000 cores), L-BFGS uses 800 (10000 cores)

©Emily Fox 2015

*Images from Duchi et al. ISMP 2012 slides*

25

## What you should know about Logistic Regression (LR) and Click Prediction

- Click prediction problem:
  - Estimate probability of clicking
  - Can be modeled as logistic regression
- Logistic regression model: Linear model
- Gradient ascent to optimize conditional likelihood
- Overfitting + regularization
- Regularized optimization
  - Convergence rates and stopping criterion
- Stochastic gradient ascent for large/streaming data
  - Convergence rates of SGD
- AdaGrad motivation, derivation, and algorithm

©Emily Fox 2015

26